

Abdul Basit Sajid

Backend / Infrastructure Engineer

basit@basit.engineer | +92 324 5717425 | Remote, based in Pakistan
basit.engineer | github.com/abd-ulbasit | linkedin.com/in/abd-ulbasit

SUMMARY

Backend and infrastructure engineer, 3 years, remote for a UK team since 2023. I work on capacity, latency and the deployment path: what a system costs to run, why it is slow, and how it ships without going down. Production is Python, Node and Go on GCP/AWS with Docker; the Kubernetes and Postgres-internals depth below is self-directed.

EXPERIENCE

Software Engineer, EliteITTeam, Manchester UK (Remote)

Sept 2024 – Present

- Moved a real-time multi-model computer-vision backend from a measured ~120-concurrent ceiling on one L4 to 500 on four L40Ses under load test at 97.9% delivery, the extra GPUs only usable once the single-process bottleneck was gone. Cut per-frame inference 2,800 ms → 60 ms (46×) and p50 871 ms → 59 ms: four YOLO models onto CUDA, an $O(n^2)$ Python NMS loop replaced with `torchvision.ops.nms`, and a `broadcast=True` fan-out removed that made every response $O(N)$ in connected sockets. Found in a 7-configuration sweep in which three of my own hypotheses were falsified.
- Replaced a stop-then-start production deploy, down for a full application boot every release, with a health-gated cutover on plain Docker and no orchestrator: the replacement container starts on a transition port and traffic only moves once it passes health checks. Still the production deploy path eight months later.
- Diagnosed why detection results silently never reached mobile clients: carrier-grade NAT placed a client's two WebSocket connections on different public IPs and therefore different containers. Shipped a 684-line Go fanout proxy with SHA-256 dedup, after four alternatives were evaluated and rejected on documented grounds.
- Wrote the capacity decision document for a 60,000-concurrent-user target: corrected an internal \$206k/mo projection to \$475–570k/mo from measured per-box capacity (90 users, not the assumed 200), established that GPUs are excluded from GCP's Flexible CUD, and recommended **not** building the GPU fleet. It was not built.
- Root-caused a Stripe webhook failure and repaired the affected production data with a dry-run pass, row checksums and a written rollback procedure before executing against live records.
- Closed external penetration-test findings on two products. Shipped IP-keyed rate limiting, then caught that all traffic arrives through one server-action IP and would bucket every user into one limiter; re-keyed it on the account.

Associate Software Engineer, EliteITTeam, Manchester UK (Remote)

Sept 2023 – Sept 2024

- Built and still maintain the backend for an edtech product: auth with device-seat licensing, Stripe billing with a webhook state machine, Apple and Google receipt verification, TTS and phoneme search.

OPEN-SOURCE INFRASTRUCTURE PROJECTS

pgoverlay: copy-on-write Postgres branching (Go, OverlayFS, Docker)

github.com/abd-ulbasit/pgoverlay

- Diagnosed why branching a 5 GiB database took 61.9s and left a 5.05 GiB writable layer. Postgres's pre-WAL-replay `SyncDataDirectory` opens every PGDATA file `O_RDWR`, and on OverlayFS an `O_RDWR` open of a lower-layer file forces whole-file copy-up.
- Proved the mechanism instead of arguing it: a single-variable control run with `recovery_init_sync_method=syncfs` ended at 16 KiB of writable layer, not 5.05 GiB. The one-flag fix takes create to 1.89s / 33 MiB.
- Branching a live branch freezes the parent's write layer into a refcounted immutable layer both then share, so you can branch a database that is still being written to. A Postgres wire-protocol router reaches every branch on one port as `dbname@branch`.

upgradescope: Kubernetes upgrade-readiness scanner (Go, client-go)

github.com/abd-ulbasit/upgradescope

- Collectors degrade instead of failing: a 401/403 from a managed control plane's metrics endpoint becomes a named unavailable capability, so the report states which signal was lost rather than reporting a clean scan.
- Argued against a controller-runtime reconciler: the verdict is a function of (inventory, knowledge base, target) and two of those three change with no Kubernetes event, so a watch-driven loop requeues on unrelated pod churn yet still misses the day a component's EOL passes.

goqueue: distributed message queue (Go)

github.com/abd-ulbasit/goqueue

- Found a bucket-drain race in a hierarchical timing wheel: an element pointer held across the callback unlock is silently invalidated by a concurrent cancel, stranding every remaining timer in that bucket for up to 7.76 days. Invisible to `-race` (every access really was under the mutex) and to the existing tests; wrote deterministic regression tests that fail against the old implementation, then fixed it.

forgepoint: ML-lifecycle platform (Go, Kubernetes, gRPC, NATS JetStream)

github.com/abd-ulbasit/forgepoint

- A closed loop from training through serving and back to retraining on drift. The engineering problem is that eleven services, each owning its own database, have to agree on what happened without a distributed transaction — including a cancelled workflow that must durably record that it was cancelled, on a context that is already dead.

WRITING, EDUCATION & CERTIFICATIONS

Bookstore Kubernetes Guide

abd-ulbasit.github.io/bookstore-kubernetes-guide

16 parts, 115 chapters: one application from a container to production EKS, through security, GitOps and day-2 ops.

Postgres copied 5 GiB before recovery started

basit.engineer/posts

The OverlayFS copy-up mechanism behind the pgoverlay diagnosis, and the control run that proved it.

BS Computer Science, NUST Islamabad | AWS Certified Solutions Architect (Associate), July 2025

SKILLS

Languages: Go, Python, TypeScript, SQL, Rust | **Data:** PostgreSQL, MongoDB, Redis | **Cloud:** AWS, GCP

Infrastructure: Kubernetes, Docker, Helm, Terraform, Nginx, GitHub Actions, Prometheus, gRPC, NATS JetStream

Practices: Distributed systems, concurrency, load testing and capacity planning, cost modelling, production migrations